

## EXP NO:2 Run a basic Word Count Map Reduce program to understand Map Reduce Paradigm.

### AIM:

To run a basic Word Count MapReduce program.

### PROCEDURE:

#### Step 1 : Create Data File:

Create a file named "word\_count\_data.txt" and populate it with text data that you wish to analyse. Login with your hadoop user.

**\$nano word\_count.txt**

### Output:

```
GNU nano 7.2 word_count.txt
Made it to LA yeah
Finally in LA yeah
Lookin for the weed though
Fryna make my own dough
Callin for Maria
Lost without Maria
Might dive in the marina
```

Help	Write Out	Where Is	Cut	Execute	Location
Exit	Read File	Replace	Paste	Justify	Go To Line

#### Step 2: Mapper Logic - mapper.py:

Create a file named "mapper.py" to implement the logic for the mapper. The mapper will read input data from STDIN, split lines into words, and output each word with its count.

**\$nano mapper.py**

### CODE:

#### Mapper.py

```
#!/usr/bin/env python3
# import sys because we need to read and write data to STDIN and STDOUT
#!/usr/bin/python3
import sys
for line in sys.stdin:
```

```
line = line.strip() # remove leading and trailing whitespace
words = line.split() # split the line into words
for word in words: nano word_count.txt
    print( '%s\t%s' % (word, 1))
```

### Step 3: Reducer Logic - reducer.py:

**\$nano reducer.py**

**CODE:**  
**Reducer.py**

```
#!/usr/bin/python3
from operator import itemgetter
import sys
current_word = None
current_count = 0
word = None
for line in sys.stdin:
    line = line.strip()
    word, count = line.split('\t', 1)
    try:
        count = int(count)
    except ValueError:
        continue
    if current_word == word:
        current_count += count
    else:
        if current_word:
            print( '%s\t%s' % (current_word, current_count)) current_count =
            count
        current_word = word
        if current_word == word:
            print( '%s\t%s' % (current_word, current_count))
```

### Step 4: Prepare Hadoop Environment:

Start the Hadoop daemons and create a directory in HDFS to store your data.

**start-all.sh**

**hdfsdfs -mkdir /word\_count\_in\_python**

**hdfsdfs -copyFromLocal /path/to/word\_count.txt/word\_count\_in\_python**

### Step 6: Make Python Files Executable:

Give executable permissions to your mapper.py and reducer.py files.

**chmod 777 mapper.py reducer.py**

### Step 7: Run Word Count using Hadoop Streaming:

Download the latest hadoop-streaming jar file and place it in a location you can easily access. Then run the Word Count program using Hadoop Streaming.

**hadoop jar /path/to/hadoop-streaming-3.3.6.jar \  
-input /word\_count\_in\_python/word\_count\_data.txt \  
-output /word\_count\_in\_python/new\_output \  
-mapper /path/to/mapper.py \  
-reducer /path/to/reducer.py**

```
hadoop@ubuntu:/home ~/wordcount$ hadoop jar /home/hadoop/Documents/hadoop-streaming-3.3.6.jar -input  
t /word_count_in_python/word_count.txt -output /word_count_in_python/output -mapper mapper.py -r  
educer reducer.py  
2023-10-25 22:54:51,391 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties  
2023-10-25 22:54:51,526 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).  
2023-10-25 22:54:51,526 INFO impl.MetricsSystemImpl: JobTracker metrics system started  
2023-10-25 22:54:51,537 WARN impl.MetricsSystemImpl: JobTracker metrics system already initialized!  
2023-10-25 22:54:51,785 INFO mapred.FileInputFormat: Total input files to process : 1  
2023-10-25 22:54:51,912 INFO mapreduce.JobSubmitter: number of splits:1  
2023-10-25 22:54:52,141 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1007878220_0001  
2023-10-25 22:54:52,141 INFO mapreduce.JobSubmitter: Executing with tokens: []  
2023-10-25 22:54:52,320 INFO mapreduce.Job: The url to track the job: http://localhost:8080/  
2023-10-25 22:54:52,322 INFO mapreduce.Job: Running job: job_local1007878220_0001  
2023-10-25 22:54:52,327 INFO mapred.LocalJobRunner: OutputCommitter set in config null  
2023-10-25 22:54:52,328 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapred.FileOutputCo  
mmitter  
2023-10-25 22:54:52,332 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2  
2023-10-25 22:54:52,333 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders  
under output directory:false, ignore cleanup failures: false  
2023-10-25 22:54:52,409 INFO mapred.LocalJobRunner: Waiting for map tasks  
2023-10-25 22:54:52,412 INFO mapred.LocalJobRunner: Starting task: attempt_local1007878220_0001_m_000000_0  
2023-10-25 22:54:52,455 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2  
2023-10-25 22:54:52,455 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders  
under output directory:false, ignore cleanup failures: false  
2023-10-25 22:54:52,493 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]  
2023-10-25 22:54:52,517 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/word_count_in_python/wo  
rd_count.txt:0+150  
2023-10-25 22:54:52,570 INFO mapred.MapTask: numReduceTasks: 1  
2023-10-25 22:54:52,647 INFO mapred.MapTask: (SPLIT) 0 len: 20234206/104057504)
```

### Step 8: Check Output:

Check the output of the Word Count program in the specified HDFS output directory.

**hdfs dfs -cat /word\_count\_in\_python/output/part-00000**

## OUTPUT:

```
hadoop@ubuntu:/home :~/wordcount$ hdfs dfs -cat /word_count_in_python/output/part-00000  
Callin 1  
Finally 1  
LA 2  
Lookin 1  
Lost 1  
Made 1  
Maria 2  
Might 1  
Tryna 1  
dive 1  
dough 1  
for 2  
in 2  
it 1
```

## RESULT:

Thus, the program for basic Word Count Map Reduce has been executed successfully.